

```

1 !pip install imblearn joblib seaborn
2
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 import logging
8 from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
9 from sklearn.preprocessing import StandardScaler, OneHotEncoder, PolynomialFeature
10 from sklearn.impute import SimpleImputer
11 from sklearn.preprocessing import StandardScaler, OneHotEncoder, PolynomialFeature
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14 from sklearn.compose import ColumnTransformer
15 from sklearn.pipeline import Pipeline
16 from sklearn.metrics import (classification_report, accuracy_score, f1_score,
17                               roc_auc_score, confusion_matrix, roc_curve,
18                               precision_recall_curve, precision_score, recall_score)
19 from sklearn.linear_model import LogisticRegression
20 from sklearn.ensemble import RandomForestClassifier, HistGradientBoostingClassifier
21 from imblearn.over_sampling import SMOTE
22 from imblearn.pipeline import Pipeline as ImbPipeline
23 import joblib
24

```

```

Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl.metadata (355 bytes)
Requirement already satisfied: joblib in /usr/local/lib/python3.12/dist-packages (1.5.2)
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.12/dist-packages (from imblearn) (0.14.0)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.0.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.12/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.12/dist-packages (from seaborn) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.0.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.10.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.22.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (20.4)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (8.3.0)
Requirement already satisfied: pyParsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.3.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.8.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: scipy<2,>=1.11.4 in /usr/local/lib/python3.12/dist-packages (from imbalanced-learn->imblearn) (1.9.3)
Requirement already satisfied: scikit-learn<2,>=1.4.2 in /usr/local/lib/python3.12/dist-packages (from imbalanced-learn->imblearn) (1.4.2)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from imbalanced-learn->imblearn) (2.2.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)
Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Installing collected packages: imblearn
Successfully installed imblearn-0.0

```

▼ File Import from PC

```

1 from google.colab import files
2
3 uploaded = files.upload()
4
5 # Get uploaded filename
6 file_name = list(uploaded.keys())[0]
7
8 # Read Excel file
9 df = pd.read_excel(file_name)
10 df.head()

```

11

Choose files osteoarthritis_dataset.xlsx
osteoarthritis_dataset.xlsx(application/vnd.openxmlformats-officedocument.spreadsheetml.sheet) - 7092587 bytes, last modified: 25/10/2025 - 100%
done
Saving osteoarthritis_dataset.xlsx to osteoarthritis_dataset.xlsx

patient_id	age	sex	bmi	knee_side	cartilage_medial_mm	cartilage_lateral_mm	joint_space_medial_mm	joint_space_lateral_mm
0	1	78	Female	32.35	Right	2.09	2.52	2.88
1	2	68	Male	28.86	Right	3.06	3.74	2.25
2	3	54	Female	27.47	Right	3.68	2.85	2.16
3	4	82	Female	22.56	Left	2.73	4.84	2.61
4	5	47	Female	31.24	Right	3.46	3.47	1.88

5 rows × 29 columns

Basic Checks

```

1 logging.basicConfig(level=logging.INFO)
2 logger = logging.getLogger(__name__)
3
4 assert df.shape[0] > 0 and df.shape[1] > 0, "Empty dataset"
5 logger.info(f"Dataset shape: {df.shape}")
6
7 print(df.shape)
8 print(df.info())
9 print(df.describe())
10
11 # Missing values
12 df.isnull().sum().sort_values(ascending=False).head()
13
14 plt.figure(figsize=(10,4))
15 df.isnull().mean().sort_values(ascending=False).plot(kind="bar")
16 plt.title("Missing Values (%) per Column")
17 plt.show()
18
19 # Class balance
20 df['progression_2yr'].value_counts().plot(kind="bar")
21 plt.title("Target Distribution")
22 plt.show()
23
24 # Numeric distributions
25 df.select_dtypes(include='number').hist(figsize=(12,10))
26 plt.show()
27
28 # Correlation heatmap
29 plt.figure(figsize=(12,8))
30 sns.heatmap(df.corr(numeric_only=True), annot=False, cmap='coolwarm')
31 plt.title("Correlation Heatmap")
32 plt.show()
33

```



```
(50000, 29)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 29 columns):
 #   Column           Non-Null Count Dtype  
--- 
 0   patient_id      50000 non-null  int64  
 1   age              50000 non-null  int64  
 2   sex              50000 non-null  object  
 3   bmi              50000 non-null  float64 
 4   knee_side        50000 non-null  object  
 5   cartilage_medial_mm 45000 non-null  float64 
 6   cartilage_lateral_mm 45000 non-null  float64 
 7   joint_space_medial_mm 50000 non-null  float64 
 8   joint_space_lateral_mm 50000 non-null  float64 
 9   osteophytes      50000 non-null  int64  
10  bone_tscore      45000 non-null  float64 
11  vms_stress_MPa   50000 non-null  float64 
12  meniscal_damage  50000 non-null  int64  
13  acl_status       50000 non-null  object  
14  pain_score       50000 non-null  int64  
15  walking_speed_m_s 50000 non-null  float64 
16  inflammation_marker_crp 45000 non-null  float64 
17  inflammation_marker_esr 45000 non-null  float64 
18  genetic_risk     50000 non-null  int64  
19  cholesterol      50000 non-null  float64 
20  hemoglobin       50000 non-null  float64 
21  socioeconomic_status 50000 non-null  object  
22  exercise_frequency 50000 non-null  int64  
23  smoking_status   50000 non-null  object  
24  alcohol_consumption 29852 non-null  object  
25  texture_entropy   45000 non-null  float64 
26  texture_contrast  45000 non-null  float64 
27  pain_catastrophizing 50000 non-null  int64  
28  progression_2yr   50000 non-null  int64  
dtypes: float64(14), int64(9), object(6)
memory usage: 11.1+ MB
None
    patient_id      age      bmi  cartilage_medial_mm \
count  50000.000000  50000.000000  50000.000000      45000.000000
mean   25000.500000  61.98650   27.987705      3.006427
std    14433.901067  12.95123   4.985826      0.800577
min    1.000000     40.00000   7.500000      0.010000
25%   12500.750000  51.00000   24.640000      2.470000
50%   25000.500000  62.00000   27.980000      3.000000
75%   37500.250000  73.00000   31.350000      3.540000
max   50000.000000  84.00000   47.500000      6.260000

    cartilage_lateral_mm  joint_space_medial_mm  joint_space_lateral_mm \
count  45000.000000      50000.000000      50000.000000
mean   3.199000          2.500796          2.799205
std    0.699206          0.598143          0.497948
min    0.440000          0.110000          0.560000
25%   2.720000          2.100000          2.460000
50%   3.200000          2.500000          2.800000
75%   3.670000          2.910000          3.140000
max   6.050000          5.250000          4.820000

    osteophytes      bone_tscore  vms_stress_MPa  ... \
count  50000.000000  45000.000000  50000.000000  ...
mean   1.504840     -1.508637  20.019953  ...
std    1.118553     1.005187  5.015887  ...
min    0.000000     -5.750000  0.040000  ...
25%   1.000000     -2.190000  16.640000  ...
50%   2.000000     -1.510000  20.000000  ...
75%   3.000000     -0.830000  23.380000  ...
max   3.000000      2.980000  42.500000  ...

    inflammation_marker_crp  inflammation_marker_esr  genetic_risk \
count  45000.000000          45000.000000  50000.000000
mean   4.991103            20.032799  0.198020
std    1.999562            10.025325  0.398511
min    -3.190000           -25.270000  0.000000
25%   3.630000            13.217500  0.000000
50%   5.000000            20.010000  0.000000
75%   6.330000            26.840000  0.000000
max   13.850000           58.880000  1.000000

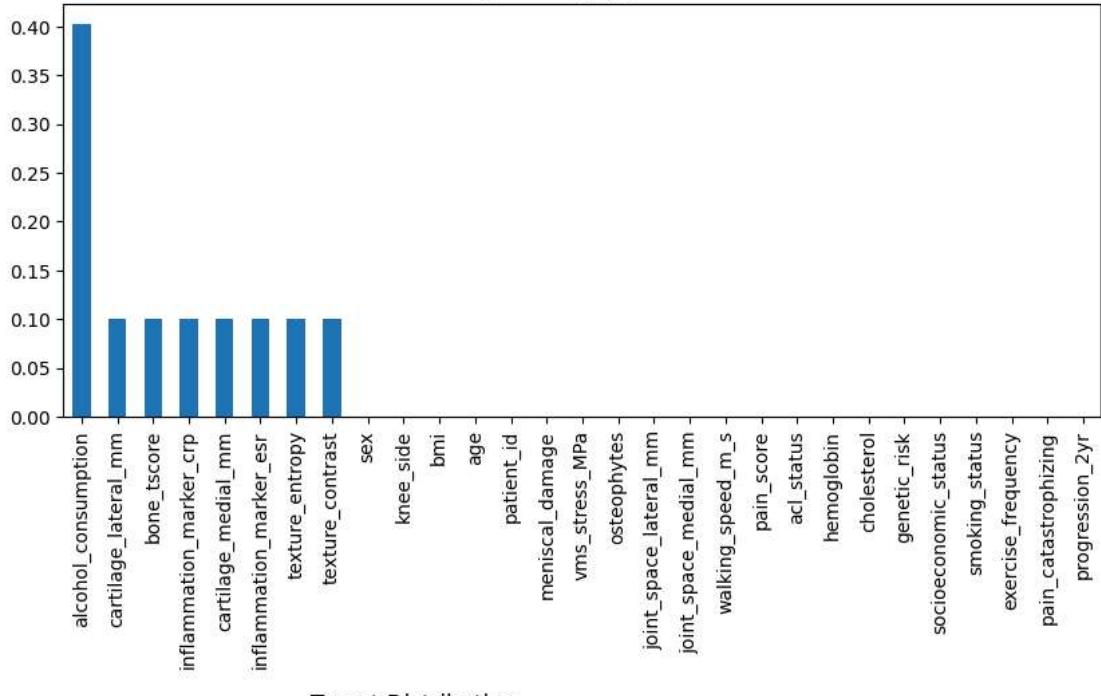
    cholesterol      hemoglobin  exercise_frequency  texture_entropy \
count  50000.000000  50000.000000  50000.000000  45000.000000
mean   199.732311  13.499771   2.498020   0.799940
std    40.105047   1.500007   1.699369   0.198822
min    14.870000   7.150000   0.000000  -0.009000
25%   172.740000  12.490000   1.000000   0.666000
50%   199.680000  13.490000   2.000000   0.801000
75%   226.740000  14.500000   4.000000   0.935000
max   372.970000  19.970000   5.000000   1.588000

    texture_contrast  pain_catastrophizing  progression_2yr
count  45000.000000      50000.000000      50000.000000
```

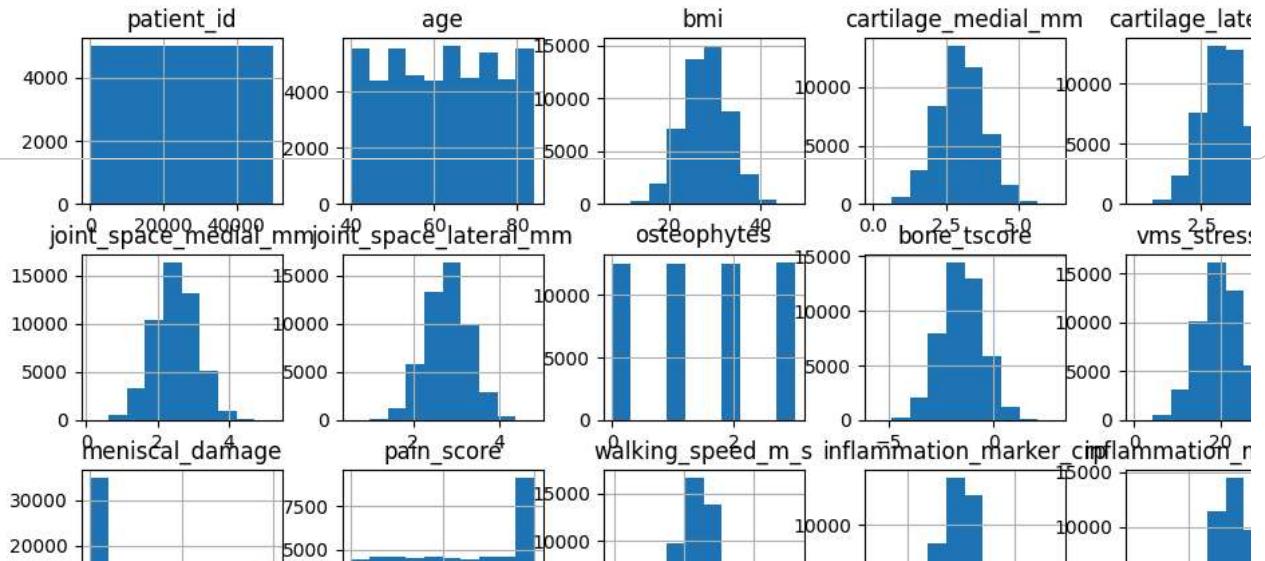
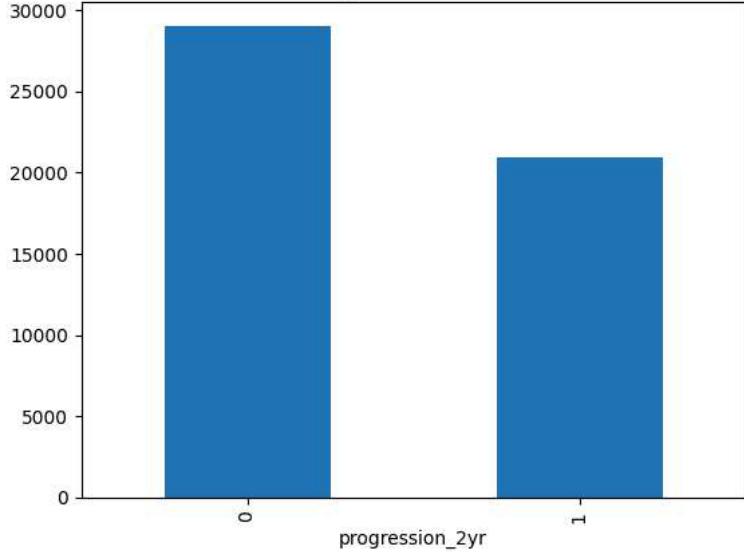
mean	1.499863	9.965480	0.419140
std	0.403010	6.069371	0.493423
min	-0.224000	0.000000	0.000000
25%	1.228000	5.000000	0.000000
50%	1.497000	10.000000	0.000000
75%	1.774000	15.000000	1.000000
max	3.332000	20.000000	1.000000

[8 rows x 23 columns]

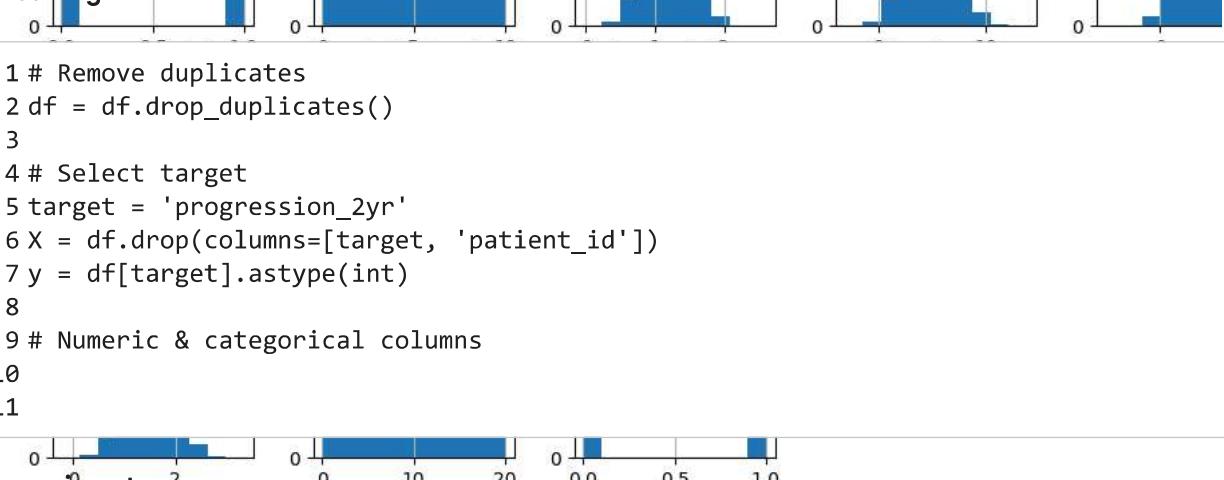
Missing Values (%) per Column



Target Distribution



▼ Data Cleaning



▼ Preprocessing step²

Correlation Heatmap

```

1 target = 'progression_2yr'
2 X = df.drop(columns=[target, 'patient_id'])
3 y = df[target].astype(int)
4
5
6 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,stratify=y,random_state=42)
7 logger.info(f'Training shape : {X_train.shape} Testing shape : {X_test.shape}')
8
9
10 numeric_cols = X.select_dtypes(include='number').columns.tolist()
11 cat_cols = X.select_dtypes(include=['object','category']).columns.tolist()
12
13 num_pipeline = Pipeline([
14     ("imputer", SimpleImputer(strategy="median")),
15     ("poly", PolynomialFeatures(degree=2, include_bias=False)),
16     ("scaler", StandardScaler())
17 ])
18
19 cat_pipeline = Pipeline([
20     ("imputer", SimpleImputer(strategy="most_frequent")),
21     ("onehot", OneHotEncoder(handle_unknown="ignore"))
22 ])
23
24 preprocessor = ColumnTransformer([
25     ("num", num_pipeline, numeric_cols),
26     ("cat", cat_pipeline, cat_cols)
27 ])
28

```

▼ Model loop

```

1 models = {
2     "LogisticRegression": {
3         "model": LogisticRegression(max_iter=2000, solver="saga"),
4         "params": {"model_C": [10]}
5     },
6     "RandomForest": {
7         "model": RandomForestClassifier(),
8         "params": {"model_n_estimators": [5,10], "model_max_depth": [1,10]}
9     },
10    "HistBoost": {
11        "model": HistGradientBoostingClassifier(),
12        "params": {"model_max_iter": [100], "model_max_depth": [10]}
13    }
14}

```

```

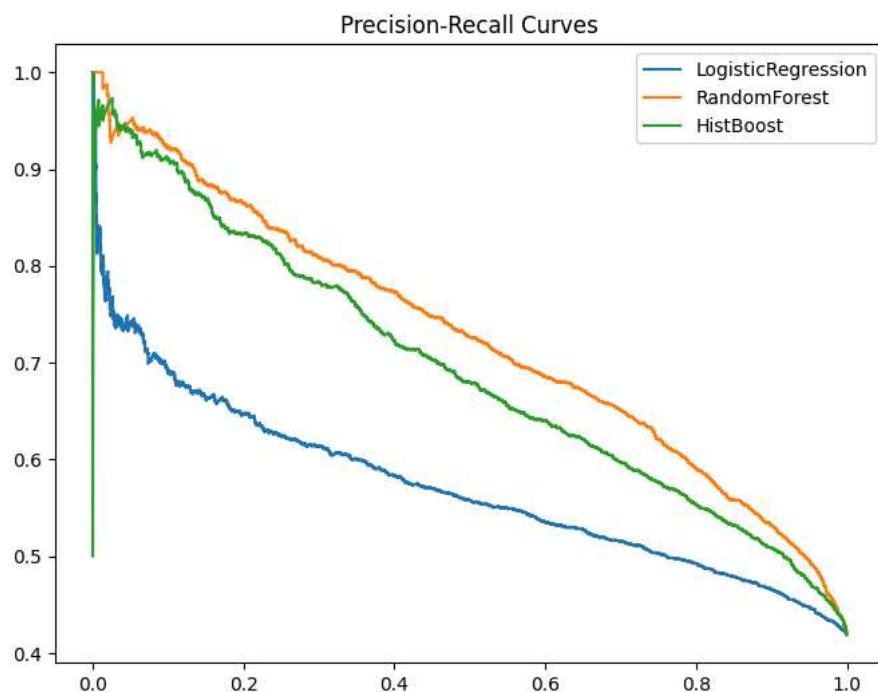
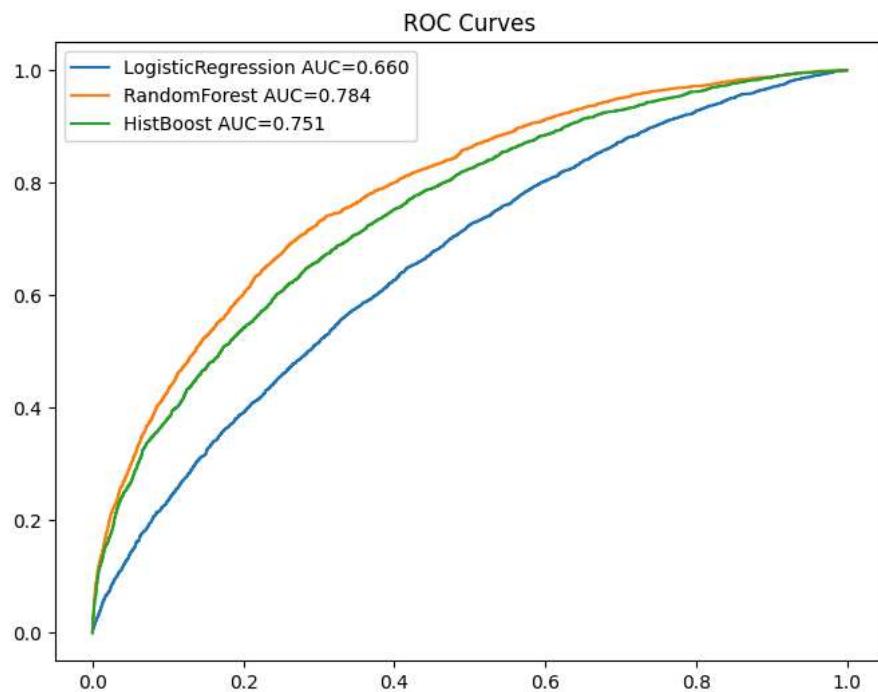
13     }
14 }
15
16 results = {}
17 skf = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)
18
19 for name, mp in models.items():
20     pipe = ImbPipeline([
21         ("preproc", preprocessor),
22         ("smote", SMOTE()),
23         ("model", mp["model"])
24     ])
25
26     grid = GridSearchCV(pipe, mp["params"], cv=skf, scoring="f1", n_jobs=-1)
27     grid.fit(X, y)
28
29     best_model = grid.best_estimator_
30     y_pred = best_model.predict(X)
31     y_proba = best_model.predict_proba(X)[:,1]
32
33     results[name] = {
34         "model": best_model,
35         "f1": f1_score(y, y_pred),
36         "roc": roc_auc_score(y, y_proba)
37     }
38
39     print(f"{name} = {results[name]}")
40
        ('smote', SMOTE()),
        ('model',
         LogisticRegression(C=10, max_iter=2000, solver='saga'))]), 'f1': 0.5788680399309727, 'roc': np.float64(0.
RandomForest = {'model': Pipeline(steps=[('preproc',
    ColumnTransformer(transformers=[('num',
        Pipeline(steps=[('imputer',
            SimpleImputer(strategy='median'))),
            ('poly',
            PolynomialFeatures(include_bias=False)),
            ('scaler',
            StandardScaler()))]),
        ['age', 'bmi',
            'cartilage_medial_mm',
            'cartilage_lateral_mm',
            'joint_space_medial_mm',
            'joint_space_lateral_mm',
            'osteophytes', 'bone_tscore',
            'vms_stres...
            'texture_contrast',
            'pain_catastrophizing']),
        ('cat',
        Pipeline(steps=[('imputer',
            'texture_contrast'
        )])
    )
}

```

```
'context',  
'pain_catastrophizing']),  
('cat',  
    Pipeline(steps=[('imputer',  
        SimpleImputer(strategy='most_frequent')),  
        ('onehot',  
            OneHotEncoder(handle_unknown='ignore'))])),  
['sex', 'knee side'].
```

ROC & PR Curves

```
1 import matplotlib.pyplot as plt  
2 plt.figure(figsize=(8,6))  
3 for name in results:  
4     model = results[name]["model"]  
5     probs = model.predict_proba(X_test)[:,1]  
6     fpr, tpr, _ = roc_curve(y_test, probs)  
7     auc_val = roc_auc_score(y_test, probs)  
8     plt.plot(fpr, tpr, label=f"{name} AUC={auc_val:.3f}")  
9 plt.legend()  
10 plt.title("ROC Curves")  
11 plt.show()  
12  
13 plt.figure(figsize=(8,6))  
14 for name in results:  
15     probs = results[name]["model"].predict_proba(X_test)[:,1]  
16     prec, rec, _ = precision_recall_curve(y_test, probs)  
17     plt.plot(rec, prec, label=name)  
18 plt.legend()  
19 plt.title("Precision-Recall Curves")  
20 plt.show()  
21
```



Model Comparison Table

```

1 comparison = pd.DataFrame(results).T
2 comparison = comparison[['acc','f1','roc','prec','rec','cv_f1_mean']]
3 print("MODEL COMPARISON:\n")
4 print(comparison)
5
6 comparison.to_csv("model_comparison.csv")
7

```

Double-click (or enter) to edit

Feature Importance (RF Only)

```

1 rf = results["RandomForest"]["model"]
2 feat_names = rf.named_steps["preproc"].get_feature_names_out()
3 importances = rf.named_steps["model"].feature_importances_

```